

Pranav Dwivedi

Gurugram, India

Phone: +91-8968250147

Email: pranavdwivedi0204@gmail.com

LinkedIn: linkedin.com/in/pranav-dwivedi-b08312221

GitHub: github.com/pranavdwivedi

Backend and AI engineer with 2+ years across full-time and internship roles, building REST APIs and LLM-powered services in Python, Django/DRF, and FastAPI on MongoDB and AWS, including production RAG pipelines and CI/CD automation.

TECHNICAL SKILLS

Languages: Python, C++

Backend: Django, Django REST Framework (DRF), FastAPI, REST API Design, Microservices, Asynchronous Programming (asyncio, aiohttp)

Databases: MongoDB (aggregation pipelines, query optimization), MySQL, Caching

AI / ML: RAG Pipelines, Vector Embeddings, Semantic Search, LangGraph, OpenAI API, HuggingFace Transformers, Pandas, NumPy

Cloud & DevOps: AWS, CI/CD (Bitbucket Pipelines), Grafana Monitoring & Alerting, Cron Scheduling

Data & Visualization: Plotly, Matplotlib, Web Scraping, Data Extraction Pipelines

Tools & Version Control: Git, Bitbucket, GitHub, Postman, Swagger, Requests, Agile workflow, Code Review

EDUCATION

I.K. Gujral Punjab Technical University

Bachelor of Technology, Computer Science and Engineering | Mohali, India | Aug 2021 – Aug 2025

EXPERIENCE

Cyber Media India Ltd. — Executive, Software Development

Gurugram, India (Hybrid) | Jan 2025 – Present

- Architected and deployed RESTful APIs for CMGalaxy using Django REST Framework, consolidating analytics from 13+ advertising platforms into a single unified data layer.
- Engineered MongoDB aggregation pipelines over high-volume marketing datasets, cutting API response time by 35%.
- Automated Google Ads and Meta Ads data extraction with Python scripts scheduled via cron jobs, replacing manual server-side pulls.
- Scripted Google Docs to WordPress conversion and publishing in Python, automating a manual content management workflow.
- Built AI-powered dashboards with Plotly and Matplotlib, surfacing ROAS, CTR, CPA, CPC, and conversion trends in one view.
- Refactored 20+ Django REST APIs from function-based to class-based views, improving maintainability and modularity.
- Implemented RAG pipelines with document vectorization, vector embeddings, and LLMs, enabling semantic search and conversational Q&A over internal files.
- Configured and maintained CI/CD pipelines in Bitbucket Pipelines, automating build, test, and deployment stages to accelerate release cycles.

Cyber Media India Ltd. — Intern, Software Development

Gurugram, India | Jul 2024 – Dec 2024

- Built an AI-driven social media content automation system in Python, covering the full pipeline from web scraping to platform-ready post generation.
- Constructed asynchronous data collection pipelines with aiohttp and asyncio, cutting ingestion time by 40%.
- Produced brand-aligned marketing content at scale using the OpenAI API and LangGraph instead of hand-written copy.
- Designed modular FastAPI components, making the content generation architecture extensible and easier to maintain.
- Reviewed pull requests and managed branches in Bitbucket within an agile team, keeping version control consistent.

PROJECTS

CMGalaxy | cmgalaxy.com

Django REST Framework, FastAPI, Python, MongoDB, Plotly, LangGraph, AWS, Grafana

- Built a centralized marketing intelligence platform, unifying ad data from 13+ sources including Google Ads, Meta Ads, LinkedIn, and Shopify.
- Shipped a RAG-based file chat feature that vectorizes user-uploaded documents, enabling semantic Q&A over them with LLMs and vector search.
- Provisioned and managed AWS services and set up Grafana, enabling live API observability, alerting, and incident tracking.
- Optimized database queries and added caching layers, reducing response latency and improving efficiency.

SmartPostAI | smartpost-ai.vercel.app

FastAPI, Python, OpenAI API, LangGraph

- Created a FastAPI backend that scrapes, cleans, and transforms web content into platform-ready social media posts and long-form articles.
- Established LLM-powered pipelines with the OpenAI API and LangGraph, automating structured creative output at scale.
- Adopted an asynchronous processing architecture, enabling concurrent content generation and higher throughput.